



Never Lose a Message Again: Building Rock-Solid Event-Driven Streaming Architectures with VMware Tanzu RabbitMQ and Spring

VMware Tanzu – Data Services



VMware Tanzu

Infrastructure for running modern apps and backing services with consistent, conformant Kubernetes everywhere.



Data Management
Management for
Tanzu Data
Services instances



GemFire

Fast In-Memory
data store for
Caching,
Transactional and
NoSQL support
powered by
Apache Geode

I need a
fast data
store



SQL

Relational MySQL
or Postgres
database for
Transactional or
Analytic data
processing

I need to
replatform a
relational
database



Greenplum

Massively Parallel
Processing (MPP)
Postgres for Big
Data store for
analytics, Machine
Learning and
Artificial Intelligence

I need to drive
analytic value
of out tons of
existing data



Rabbit MQ

High throughput broker for
reliable messaging delivery

I need reliable
messaging delivery



Spring Cloud Data Flow

Data integration
orchestration service for
dynamically building data
pipelines

I need flexible and
manageable data
integrations

Features

- ✓ Cloud deployed backing-services
- ✓ On-Premise and Multi-Cloud
- ✓ Self – Service
- ✓ Scaling
- ✓ HA - Fault Tolerant
- ✓ Based on open source
- ✓ World Class Support

RabbitMQ – 101 – Broker, Producers & Consumers

RabbitMQ is a message broker

- stores and forwards binary blobs of data – messages.

Producer

- Program that sends messages is a producer

Consumer

- Program that mostly waits to receive messages:



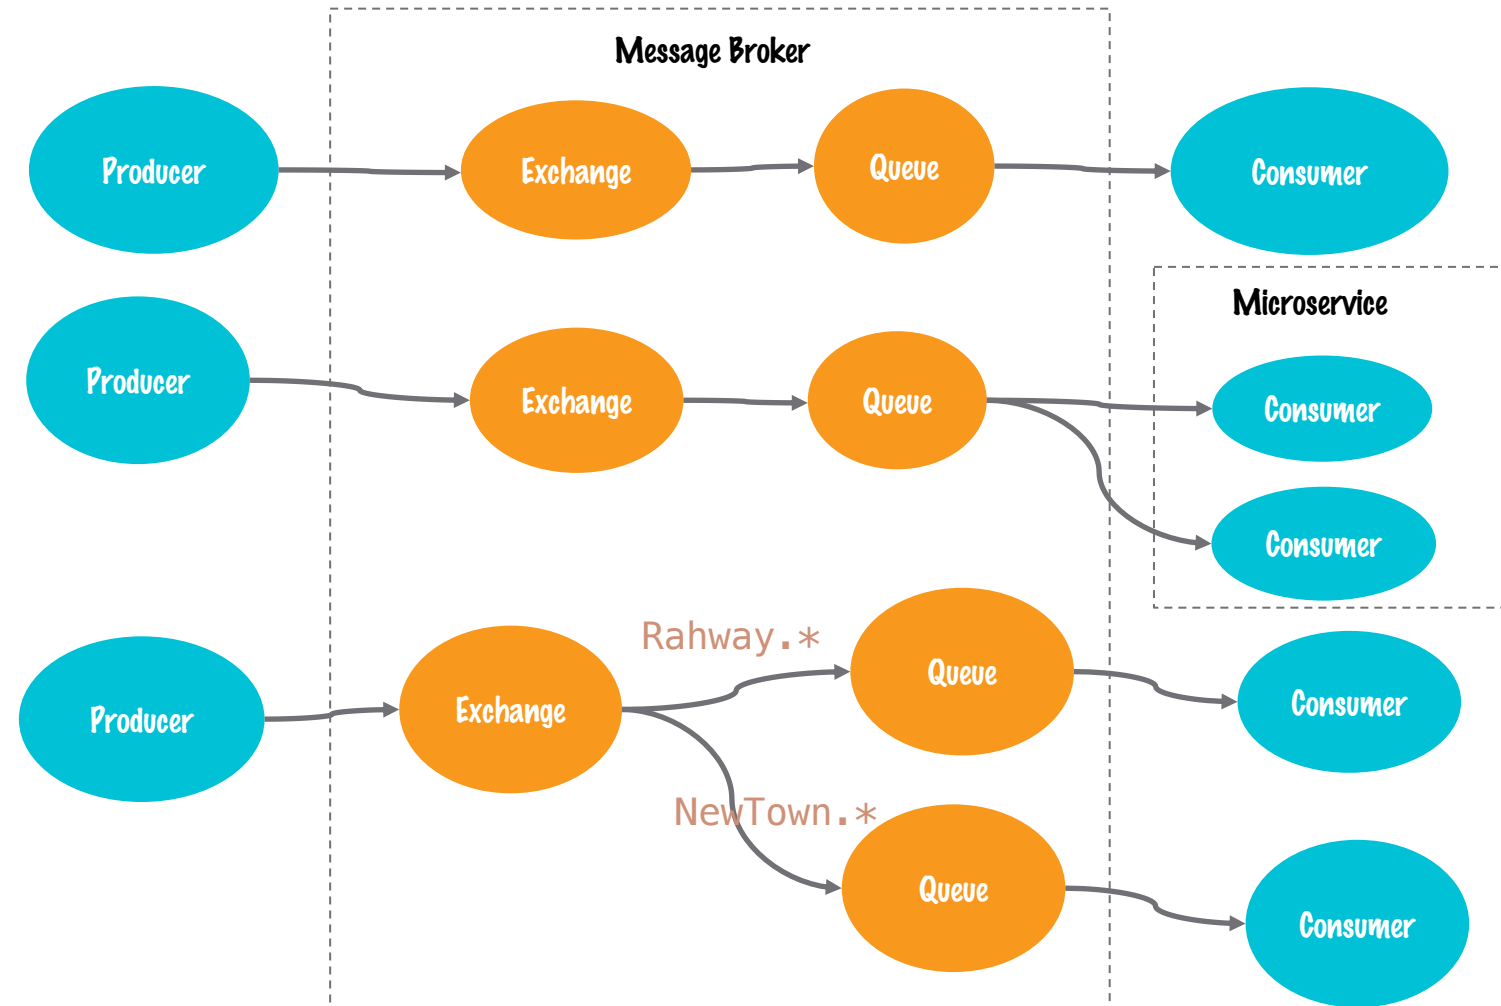
RabbitMQ – Exchanges & Queues

Queue

- Storage destination of messages inside RabbitMQ

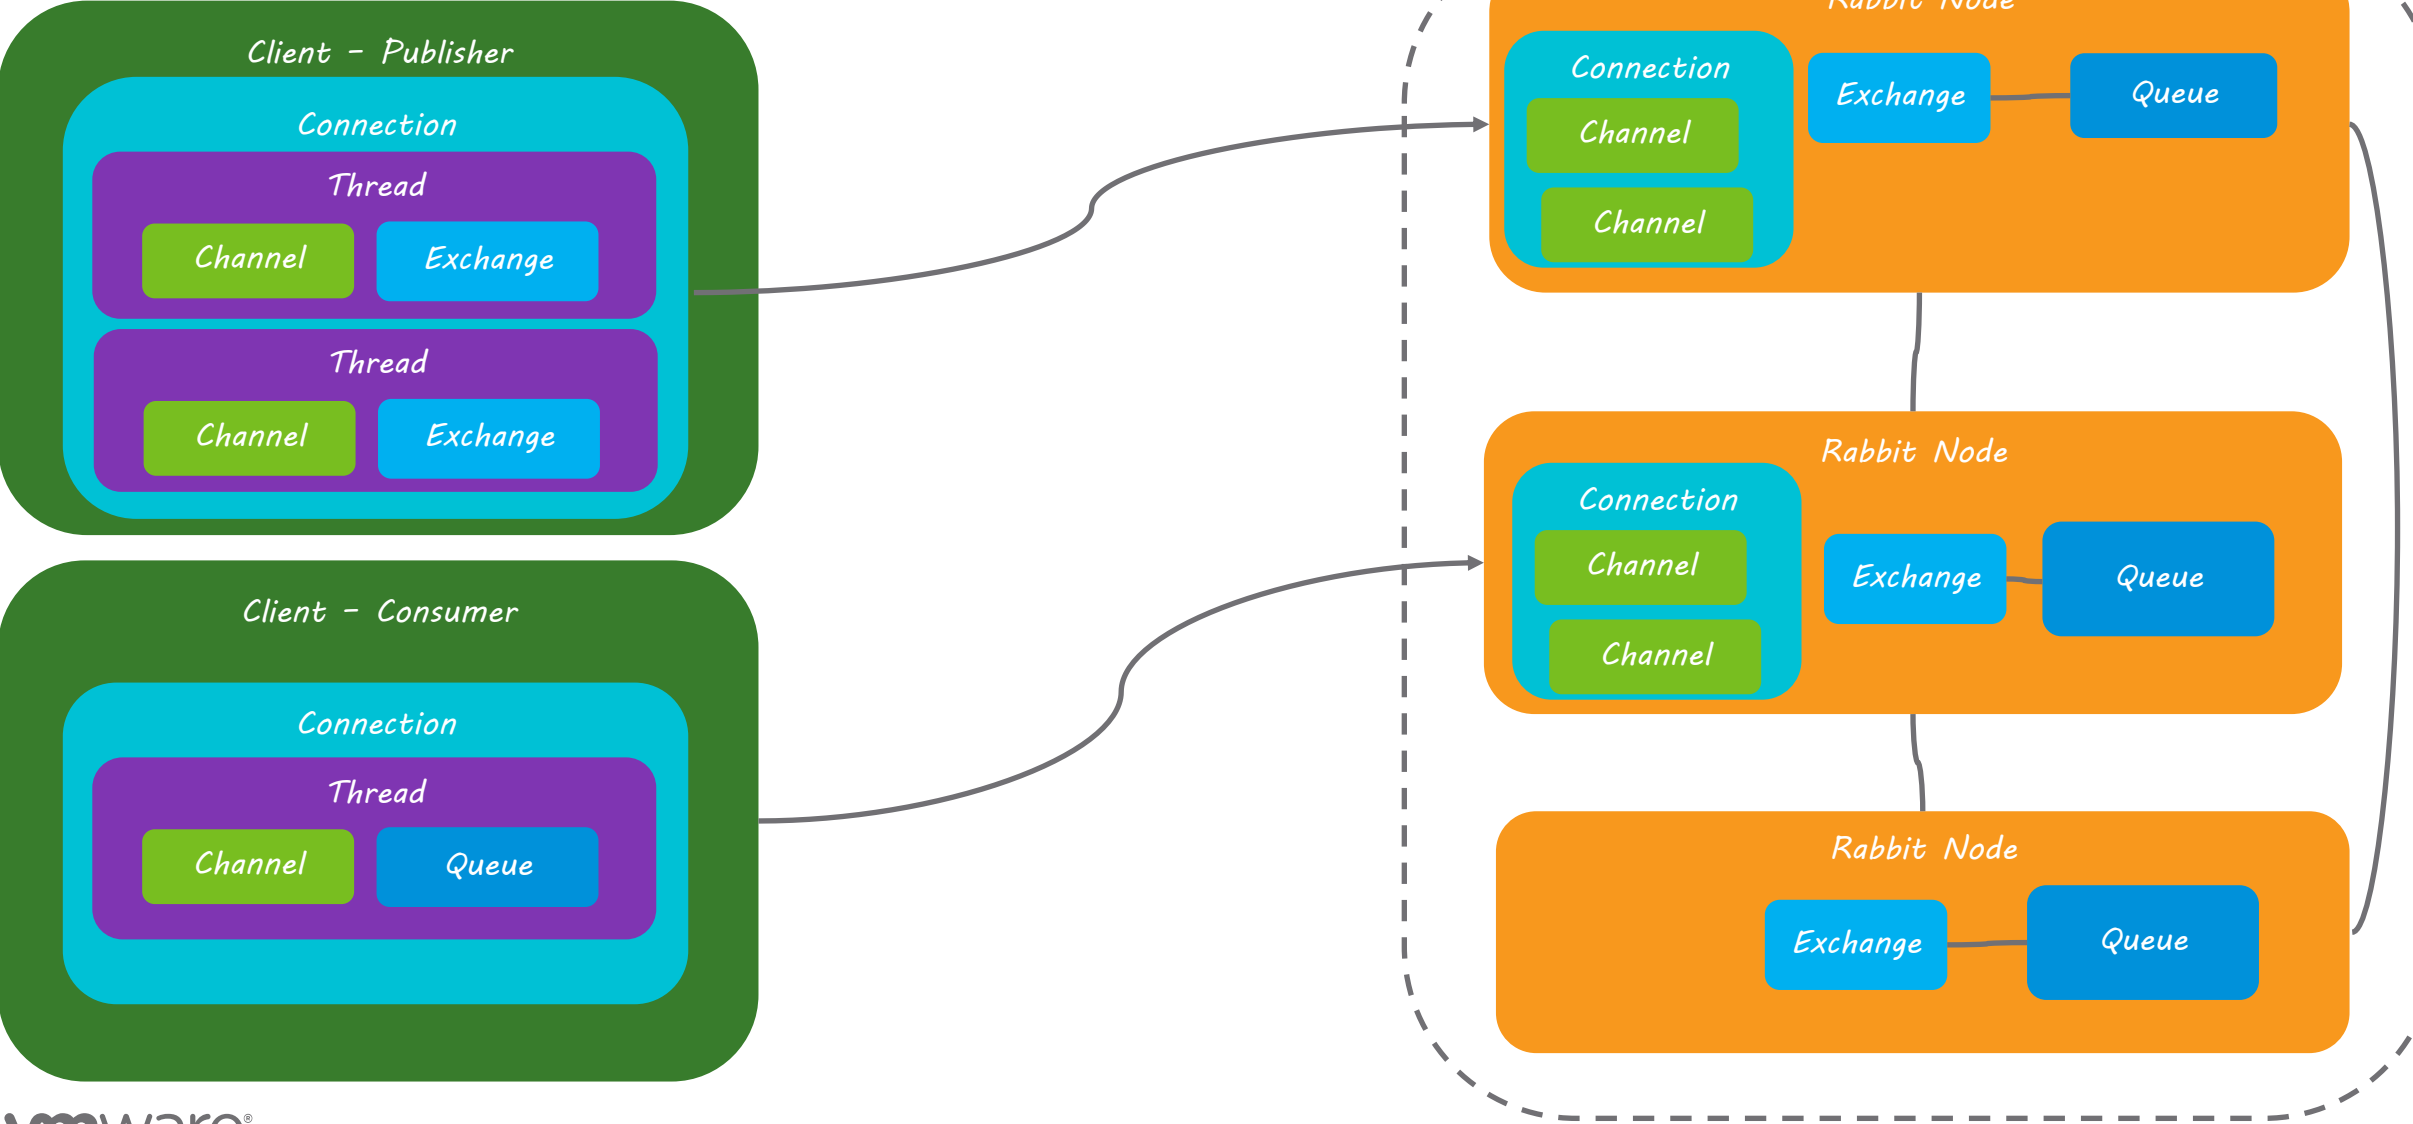
Exchanges

- Messages always sent to exchanges, then forwarded to queues based on routing rules.



RabbitMQ Clustering

Connection management



Queues

Classic versus Quorum

- Classic Queues
 - Supports In-Memory messages
 - Option durable and or persisted messages
 - Mirrored replication through policy (deprecated)
- Quorum Queues
 - A durable, replicated with persisted messages
 - Based on the Raft consensus algorithm.
 - Preferred queue type over durable mirrored classics queues.
 - Quorum queues should be considered the default option for a replicated queue type.

The screenshot displays the RabbitMQ Admin interface. At the top, there are tabs: Overview, Connections, Channels, Exchanges, Queues, and Admin. The 'Queues' tab is selected. Below the tabs, there's a 'Definition:' section with a table of parameters for a queue. The parameters are: ha-mode (exactly), ha-params (2), ha-sync-mode (automatic), ha-promote-on-shutdown (when-synced), and ha-promote-on-failure (when-synced). Below this, there are sections for 'Queues [All types]', 'Queues [Classic]', 'Queues [Quorum]', 'Exchanges', and 'Federation'. The 'Add / update policy' button is visible. Below the policy section, there's a 'Add a new queue' section. In this section, the 'Type' is set to 'Quorum'. The 'Name' is 'quorum'. The 'Arguments' section shows 'x-max-length' set to '100000000'. Below the arguments, there are various configuration options: 'Auto expire', 'Overflow behaviour', 'Single active consumer', 'Dead letter exchange', 'Dead letter routing key', 'Max length', 'Max length bytes', 'Delivery limit', 'Max in memory length', 'Max in memory bytes', and 'Initial cluster size'.

Parameter	Value	Type
ha-mode	exactly	String
ha-params	2	Number
ha-sync-mode	automatic	String
ha-promote-on-shutdown	when-synced	String
ha-promote-on-failure	when-synced	String

Queues [All types] Max length | Max length bytes | Overflow behaviour | Dead letter exchange | Dead letter routing key

Queues [Classic] HA mode | HA params | HA sync mode | HA mirror promotion on shutdown | HA mirror promotion on failure | Message TTL | Auto expire | Lazy mode | Master Locator

Queues [Quorum] Max in memory length | Max in memory bytes | Delivery limit

Exchanges Alternate exchange

Federation Federation upstream set | Federation upstream

Add / update policy

▼ Add a new queue

Type: Quorum

Name: quorum

Arguments: x-max-length = 100000000 (Number) | (String)

Add Auto expire | Overflow behaviour | Single active consumer | Dead letter exchange | Dead letter routing key | Max length | Max length bytes | Delivery limit | Max in memory length | Max in memory bytes | Initial cluster size

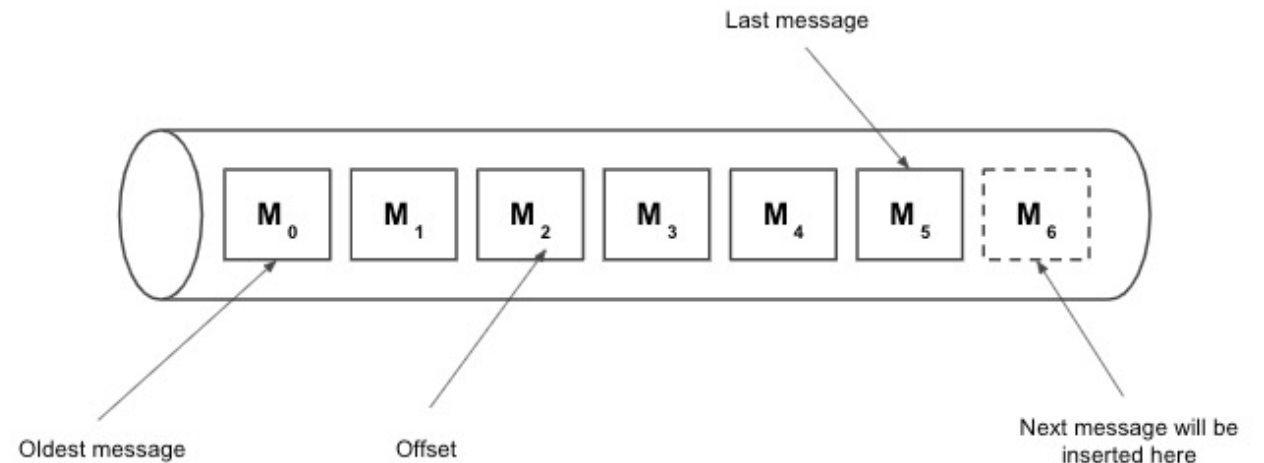
Add queue

Messaging Streaming - Queue

Replay Messages

- Kafka like event logging
- **Large fan-outs:** when several consumer applications need to read the same messages.
- **Replay / Time-traveling:** when consumer applications need to read the whole history of data or from a given point in a stream.
- **Throughput performance:** when higher throughput than with other protocols (AMQP, STOMP, MQTT) is required.
- **Large logs:** when large amount of data need to be stored, with minimal in-memory overhead.

```
channel.queueDeclare(  
    "my-stream",  
    true,          // durable  
    false, false, // not exclusive, not auto-delete  
    Collections.singletonMap("x-queue-type", "stream")  
);
```

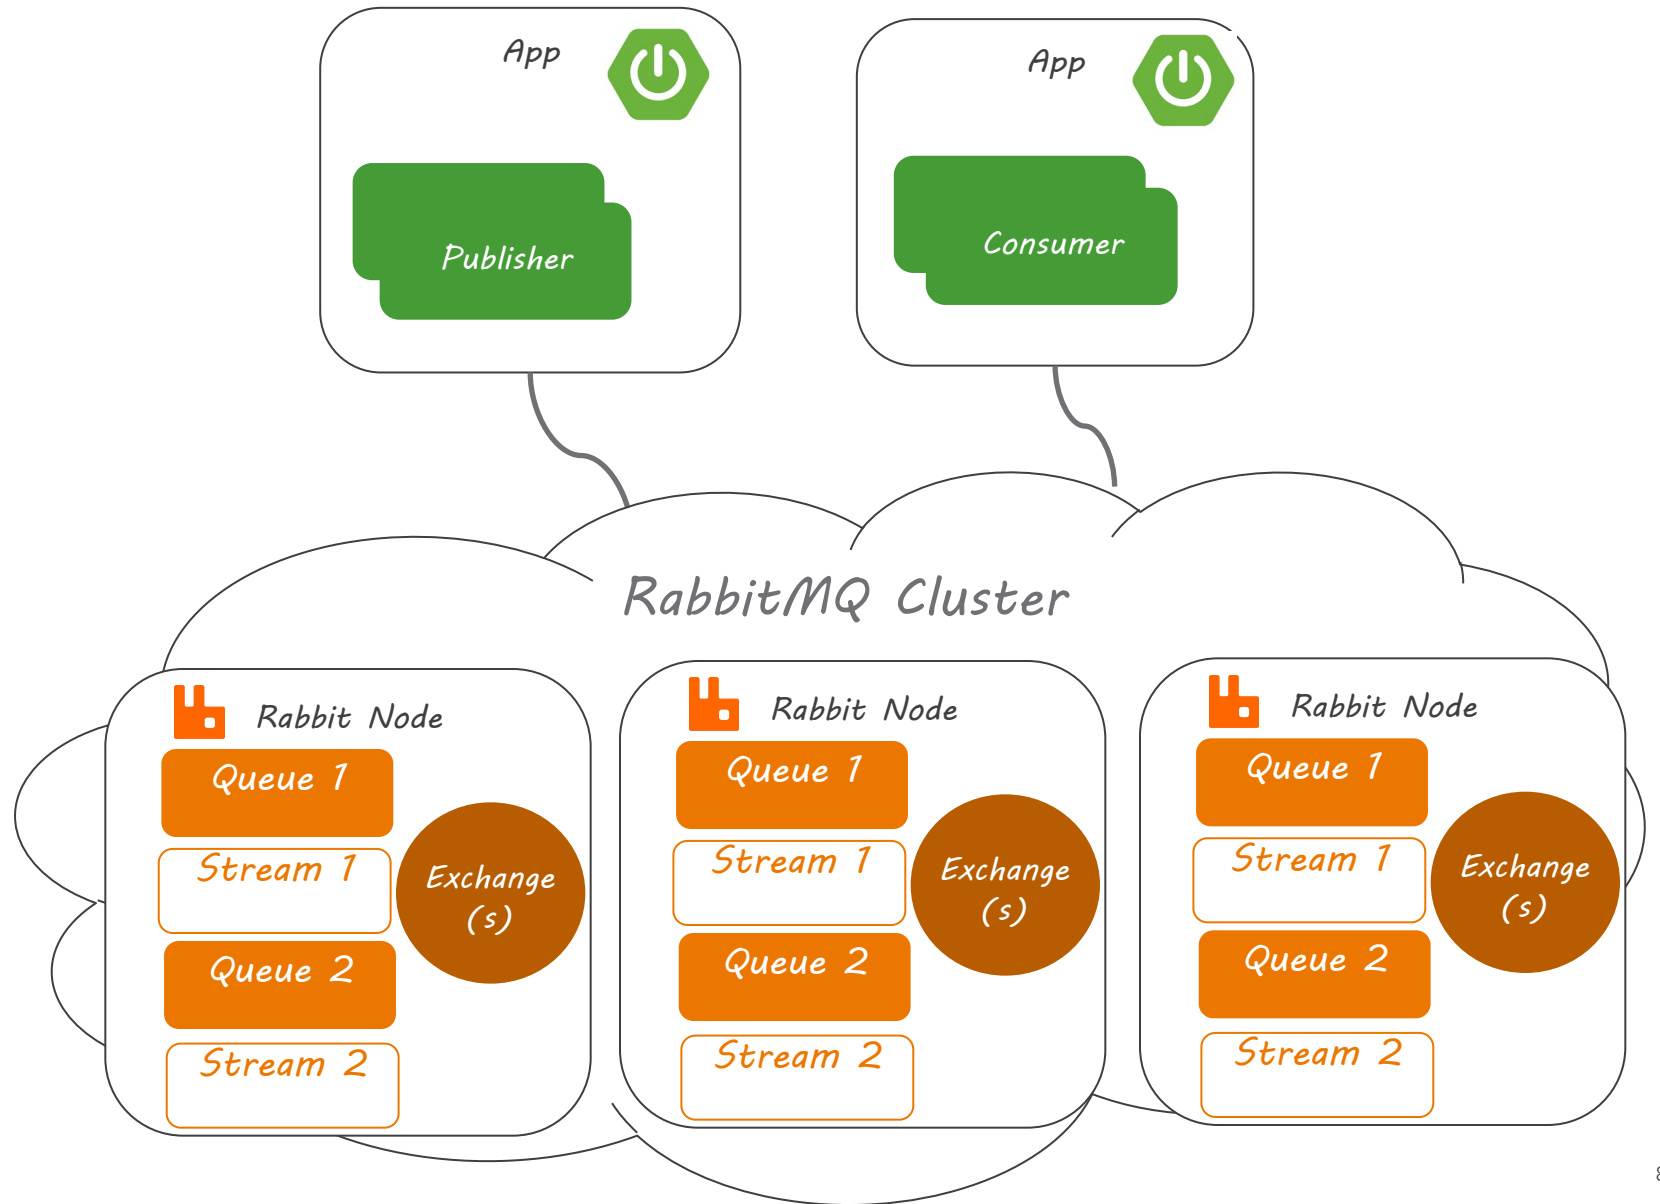


```
channel.basicConsume(  
    "my-stream",  
    false, // not auto-ack  
    Collections.singletonMap("x-stream-offset", 0),  
    (s, delivery) -> { }, // delivery callback  
    s -> { } // cancel callback  
);
```

RabbitMQ

Scalability/Reliability

- An odd number of cluster nodes are recommended (ex: 1, 3, 5, 7) by several features that require a consensus between cluster member
- A client can connect to any node.
- Nodes will route operations to the leader.
- Add more queues for Scalability across the cluster



Spring AMPQ

Publisher

- Define a Topic Exchange bean for automatic exchange creation
- RabbitTemplate can be used for sending messages
- Use @Transactional for Publisher confirms

```
@RestController("/obp/v4.0.0")
class AccountPublisherController(...) {

    init {...}

    @PostMapping("banks/{bankId}/accounts")
    @Transactional
    fun createAccount(@PathVariable("bankId") bankId: String,
                     @RequestBody account: Account): ResponseEntity<Account> {
        rabbitTemplate.convertAndSend(exchangeId, bankId, account)

        return ResponseEntity.ok(account);
    }
}
```

```
@Bean
fun exchange(@Value("\${spring.cloud.stream.bindings.supplier-out-0.destination:banking-account}")
            exchangeName: String) : Exchange
{
    return TopicExchange(exchangeName)
}
```

Spring Cloud Stream

Publishers

- Publisher
implement `java.util.function.Supplier`
- `spring.rabbitmq.publisher-confirm-type`
 - SIMPLE
 - Use `RabbitTemplate#waitForConfirms()` (or `waitForConfirmsOrDie()`) within scoped operations.

```
@Component
class AccountGeneratorSupplier(...) : Supplier<Account> {

    override fun get(): Account {
        var account = nextAccount()
        log.info( message: "account: account {}", account)
        return account
    }
}
```

```
spring:
  rabbitmq:
    publisher-confirm-type: simple
```

Spring Cloud Stream

Consumers

- Publisher implement `java.util.function.Consumer`
- Default AcknowledgeMode = AUTO
 - Auto - the container will issue the ack/nack based on whether the listener returns normally, or throws an exception.
 - `spring.cloud.stream.rabbit.bindings.<channelName>.consumer..`

```
@Component
class AccountConsumer(private val accountService: AccountService) :
    Consumer<Account> {
    override fun accept(account: Account) {
        accountService.createAccount(account)
    }
}
```

Spring Cloud Data Flow

Build, Deploy, and Monitor streaming and batch data pipelines

- Spring Cloud Data Flow for VMware Tanzu automates the deployment of data pipelines backed by cloud native applications

Spring Cloud Stream

- Spring Cloud Stream is a framework for building highly scalable event-driven microservices connected with shared messaging systems.

Dashboard

- GUI for managing data pipelines

DSL

- Pipeline definitions language similar to UNIX commands
 - Ex: file | s3

REST API & shell interface

```
dataflow:>stream list
```

The screenshot shows the 'Data Flow Pro' dashboard. The sidebar on the left contains navigation links: Applications, Streams (selected), Runtime, Tasks / Jobs, Manage, Audit Records, and Tools. The main content area is titled 'Create stream(s)'. It features a text editor with a DSL definition for a stream named 'vehicle-generator-source'. The DSL code is as follows:

```
1 vehicle-generator-source --server.port=8080 --vehicleCount=10 --messageCount=10000 --distanceIncrements=1 --delayMs=5 --  
vinPrefix='FLEET-A' | vehicles-geode-sink --spring.cloud.stream.bindings.vehicleGemFireSink-in-0.consumer.concurrency=5  
--server.port=8080 --spring.data.gemfire.pool.locators=gemfire1-locator-0.gemfire1-locator[10334]
```

Below the DSL editor, there is a visual diagram showing the connection between the source and sink components. The source component is labeled 'vehicle-generator-source' and the sink component is labeled 'vehicles-geode-sink'. The diagram shows a flow from the source to the sink. At the bottom of the interface, there are buttons for 'CANCEL' and 'CREATE STREAM(S)'.

Exercises

See <https://github.com/Tanzu-Solutions-Engineering/tanzu-rabbitmq-event-streaming-showcase.git>

- Lab 1 - Setup RabbitMQ on K8
- Lab 2 - Create a RabbitMQ Cluster with HA
- Lab 3 - Spring Apps with Quorum Queues
- Lab 4 - Spring Apps with Streams
- Lab 5 - Spring Cloud DataFlow
- Lab 6 - Provision RabbitMQ Topology Operation
 - Users, Permissions, Queues, Vhost, etc.